

Gist — Product Spec: Anti-Slop Browser Extension

Sep 29, 2026 · @Ashwin Kumar

Summary

Gist is a browser extension that dims low-value content on Google, YouTube, Pinterest and Amazon, and pulls the useful part out of pages you still need. It rates pages on how much they tell you, not on whether AI wrote them.

Problem. Search results and feeds are filling with mass-produced filler: SEO articles, AI images of things that don't exist, and product roundups where nothing was tested. People waste time clicking through pages that restate the same shallow answer.

Target users

- Researchers, students and developers tired of SEO-bait tutorials
- Home cooks and hobbyists, whose searches (recipes, crafts, DIY) are among the most polluted
- Artists, designers and shoppers hit by AI images on Google Images and Pinterest
- Privacy-minded users who already run ad blockers or use alternative search

Name. "Gist" is a working title, chosen because it matches the core promise: get the gist, skip the filler. Domain and trademark checks are still pending.

Landscape and positioning

The space is active but unowned: every existing tool is either locked to one platform, very small, or only available inside a paid search engine.

Tool	Where it works	Approach	Limitation
Kagi SlopStop	Kagi Search only	Community flags checked against Kagi's own signals; shows a slop score and downranks domains that mostly publish AI content (help page)	Only helps paying Kagi users
Slop Filter	X, LinkedIn, Reddit, YouTube comments	Scores the cadence of each post and collapses likely AI writing	Social posts only; judges style, not value
AI Slop Filter	LinkedIn	Hides posts with obvious tells such as em-dashes	About 128 users; crude signals
SlopBlock	YouTube	Crowd flags with a trust threshold; warning icons on thumbnails	One platform
Slop Evader	Web search	Only returns content published before ChatGPT's release (Nov 30, 2022)	Throws away all new content, good or bad

The gap. No one offers a mainstream, cross-platform filter that works where people already search and shop. Most tools also ask "was this written by AI?", which gets harder every model release and is unfair to human writers.

Positioning. Gist is Kagi-level slop protection for people who use Google, YouTube, Pinterest and Amazon, and it judges usefulness rather than authorship.

Core concept: a nutrition label for content

Gist asks one question of every page: does it tell me anything I couldn't get elsewhere? A human-written SEO page that says nothing scores low. A useful AI-assisted page scores fine.

This framing has three advantages over AI detection:

- **It stays accurate as models improve.** Filler is still filler, however fluent it sounds.

- **It is fair.** Non-native English speakers and formulaic-but-honest writers are not punished for style.
- **It is defensible.** Gist rates usefulness, so it never accuses a site of using AI.

The label. Hovering over any search result opens a small card. It shows an overall grade from 0 to 100, a one-word verdict (Solid, OK, Thin, Filler or Slop), and a bar for each of the five scoring dimensions below. A "Why?" link lists the specific signals behind the grade.

Scoring rubric

Each page gets a 0–100 score: five dimensions, each scored 0–100, combined by weight. Weights are starting values to be tuned against labeled data.

Dimension	Weight	What it measures	Example signals	Scored by
Information value	30%	Does the page actually answer the query?	Words before the first useful sentence; ratio of answer to total length; count of concrete specifics (numbers, units, names, dates, steps)	Local + deep scan
Originality	25%	Does it add anything beyond other top results?	Semantic overlap with the top 10 results; facts found nowhere else; original images vs stock or AI images	Deep scan
Human presence	20%	Is there evidence a real person with experience made it?	Named author with a posting history over time; first-person detail that can be checked; varied original photos; real comments	Local + domain

Dimension	Weight	What it measures	Example signals	Scored by
Site behavior	15%	Does the domain act like a content farm?	Articles per day; topic sprawl; content volume vs domain age; output spike after 2022; expired-domain takeover	Domain list
Monetization pressure	10%	Is the page built for ads more than readers?	Ad slots per 1,000 words; affiliate link density; popups before content	Local

Style tells are a tie-breaker only. Stock AI phrases and uniform sentence rhythm can move a score by at most 5 points. They are weak evidence and punish honest writers.

Score bands and what the user sees

Score	Verdict	What Gist does
80–100	Solid	Nothing, or an optional green dot
60–79	OK	Nothing
40–59	Thin	Small "thin" tag; full label on hover
20–39	Filler	Result dimmed; "Cut to the answer" offered
0–19	Slop	Result collapsed with its reason; one click expands it

Guardrails

- **Low confidence never goes below Thin.** If only local signals are available, the result can be tagged but never dimmed or collapsed.
- **Confirmed farms are capped.** A domain confirmed by community flags and review scores no higher than 25.
- **Verified humans get a floor.** A domain on the verified-human list scores at least 50, unless a page-level deep scan disagrees.

Features


The MVP targets Google search only, where slop hurts most and the page layout is predictable. Other platforms follow once scoring is trusted.

Feature	What it does	Phase
Result dimming	Applies the score bands to Google results	MVP
Nutrition label	Hover card with grade, verdict and five dimension bars	MVP
"Why was this dimmed?"	Lists the signals behind every flag	MVP
Flag button	One-click "slop" or "fine" feedback on any result	MVP
Cut to the answer	Pins the one useful fact to the top of recipe and how-to pages	MVP
Show me the humans	Sidebar of forum threads, established creators and small blogs for the same query	v0.2
Weekly slop report	Shareable card: share of results that were filler, time saved	v0.2
Subscribable filter lists	Topic lists (cooking farms, fake tech tutorials) run by named curators	v0.2
Slop-free image mode	Hides AI images on Google Images and Pinterest	v1
YouTube support	Flags faceless AI channels and AI-voiced videos	v1
Shopping protection	Flags untested "best of" roundups, suspicious review clusters, AI-written books on Amazon	v1
Firefox and Safari	Ports beyond Chrome	v1

Detection architecture

Most scoring happens on the device; the server only sees search result URLs and user flags.

Only result URLs and flags ever leave the device



detection architecture · 3 layers across device and server

Each result passes through the local layers first, and only uncached URLs trigger a paid deep scan. Flags flow into review, which updates the domain list every device downloads.

Privacy commitments

- No browsing history leaves the device, only the URLs of search results being scored.
- The Layer 1 code is open source so anyone can verify what runs locally.
- Flags are tied to an account for trust weighting, never to browsing data.

Community and trust

Community flags power the domain list, so every flag is weighted by the flagger's track record and no single user can sink a site.

- **Flag reasons.** Users pick one: filler, AI images, fake reviews, untested roundup, or other. Reasons feed the "Why?" panel.
- **Trust weighting.** A new account's flag counts for 0.2 of a vote. Weight rises as the user's flags agree with reviewed outcomes.
- **Domain threshold.** Domain-level action needs flags from at least 10 independent trusted users plus agreement from automated site-behavior signals.
- **Anti-brigading.** Bursts of flags on one domain from new accounts or shared networks are held for review, not counted. This matters because SEO competitors will try to flag each other.
- **Appeals.** Site owners prove ownership with a DNS TXT record and request review. A human decides within 7 days, and the reasons are published.
- **Curated lists.** Topic lists have named maintainers. Users choose which lists to subscribe to, as with uBlock Origin.

Go-to-market and business model

Growth comes from content people want to share; revenue comes from a small Pro tier and, later, licensing the domain reputation data.

Growth loops

- **Before/after screenshots** of the same Google search, with and without Gist. This is the launch asset.
- **Weekly slop report cards** that users share on social media.
- **A monthly public Slop Index:** worst domains, most polluted topics, platforms getting worse. It gives journalists a recurring story.
- **Creator evangelists:** independent recipe bloggers, artists and reviewers whose work Gist helps surface.
- **Launch channels:** Hacker News, relevant Reddit communities, Product Hunt.

Pricing (USD)

Tier	Price	Includes
Free	\$0	Google dimming, nutrition label, flagging, domain list

Tier	Price	Includes
Pro	\$4/month or \$36/year	Cut to the answer, deep scans, image mode, all platforms, custom lists
Data license	Custom	Domain reputation API and dataset for business buyers

B2B buyers. Advertisers already pay brand-safety vendors to avoid "made-for-advertising" junk sites, and Gist's domain data maps onto that problem. Schools and libraries (clean-source lists) and AI companies (clean training and retrieval sources) are secondary buyers.

Risks and mitigations

The biggest risk is false positives: one wrongly dimmed legitimate site damages trust more than ten missed slop pages.

Risk	Mitigation
False positives anger users and site owners	Low-confidence guardrail, visible reasons on every flag, 7-day appeals
Google or Kagi builds this natively	Google earns ads on many slop pages, so has mixed incentives; Kagi only serves its own users. Moat is cross-platform reach, community lists and the dataset
Platform layout changes break the extension	Selector config served from the backend so fixes ship without a store review; automated daily layout checks
Content farms adapt (fake anecdotes, fake authors)	Value-based scoring is harder to game than style detection; domain behavior signals are costly to fake
Deep-scan costs grow with users	Score each URL once and share the cache; most slop sits on a small set of domains
Privacy concerns block adoption	No browsing history leaves the device; local layer open-sourced

Risk	Mitigation
Chrome Web Store policy changes	Keep permissions minimal; ship Firefox early as a fallback

30-day MVP plan

Ship a Google-only Chrome extension in four weeks, then judge it on retention and false-positive rate before expanding.

1. **Week 1:** Chrome extension on Google results. Local heuristics plus a seed list of a few hundred known content farms. Result dimming and the nutrition label.
2. **Week 2:** Flag button, shared backend, per-URL score cache, "Why was this dimmed?" panel.
3. **Week 3:** Cut to the answer for recipe and how-to pages. Deep scan for Pro testers.
4. **Week 4:** Launch on Hacker News, Reddit and Product Hunt with before/after screenshots and the first Slop Index.

Success metrics (targets at day 60)

Metric	Target
Installs	5,000
Week-4 retention	35% or higher
Upheld appeals as share of dimmed domains	Under 5%
Users who flag at least once	20%
Free-to-Pro conversion	2%

Targets are starting guesses, to be revised after the first two weeks of data.

Open questions

- ☐ Is "Gist" available as a domain and trademark?
- ☐ Should the launch wedge be Google search or Pinterest images?
- ☐ Is 10 trusted users the right domain threshold?
- ☐ Should Pro offer regional pricing (for example, in INR)?

- ☐ Which LLM runs deep scans, and what does each scan cost?